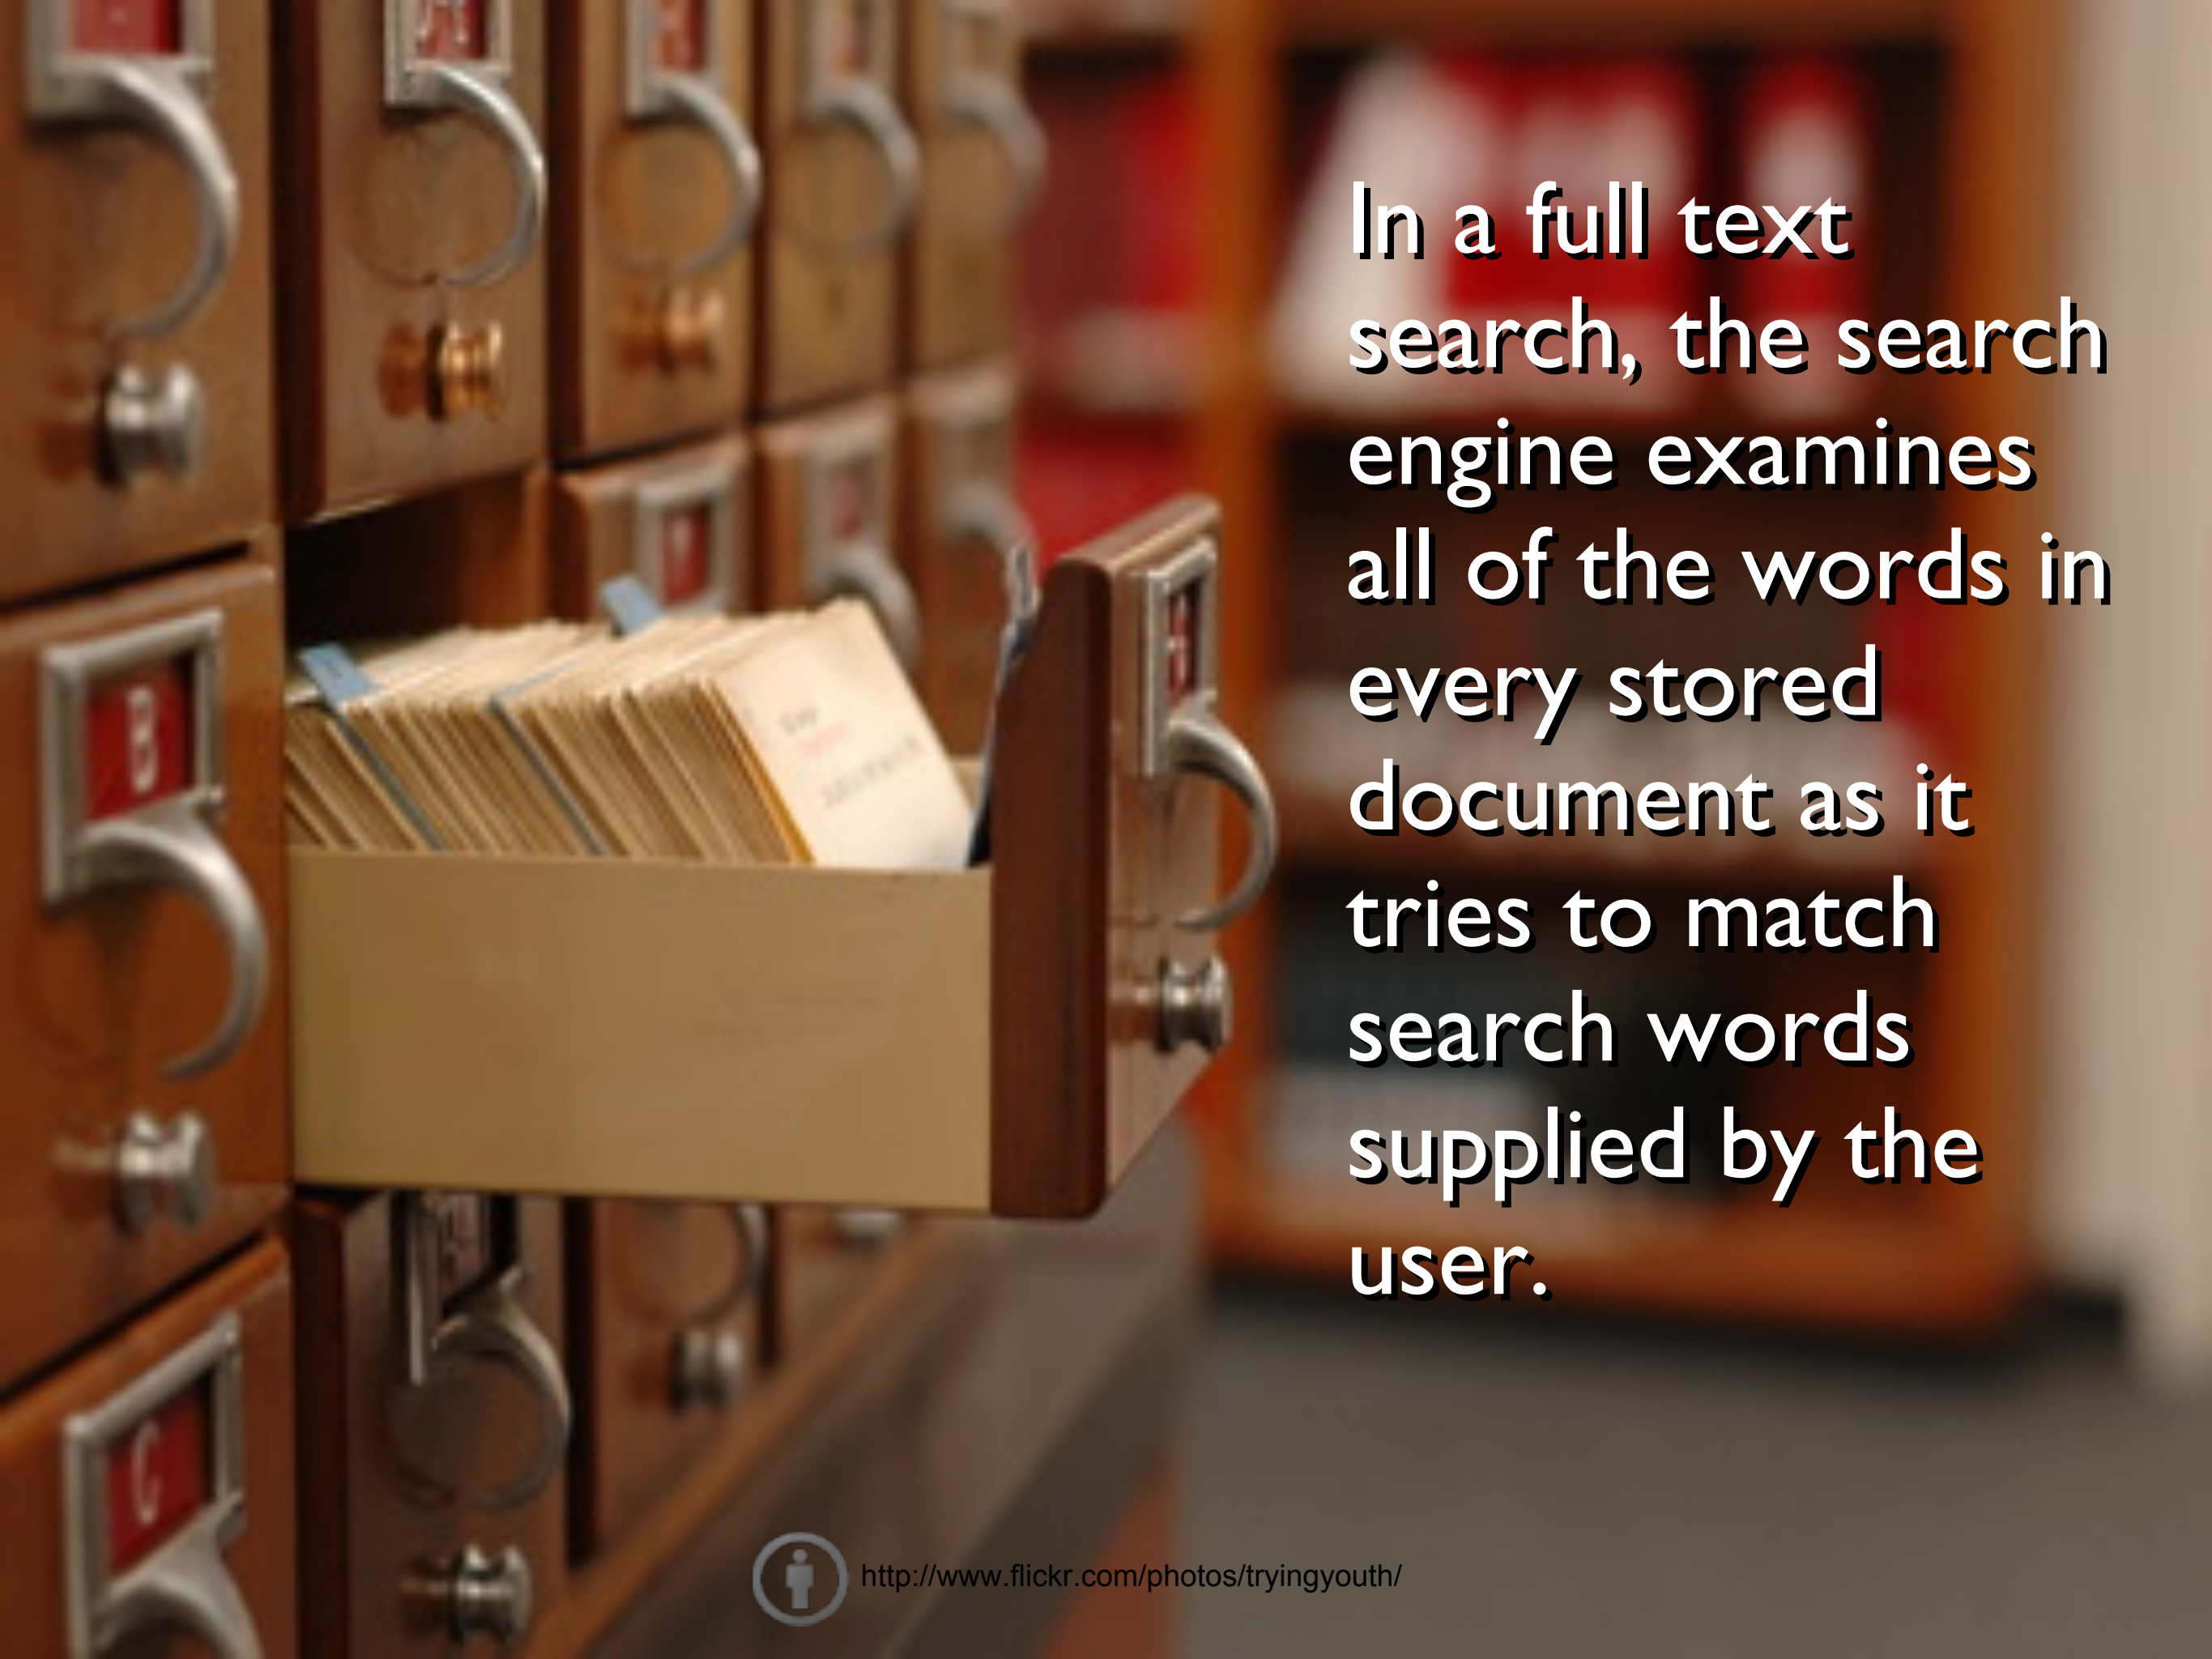




Full Text Search Throwdown

Bill Karwin, Percona Inc.



In a full text search, the search engine examines all of the words in every stored document as it tries to match search words supplied by the user.




StackOverflow Test Data

- Latest data dump, exported 2014
- 8 million Posts = 8.9 GB

Bill Karwin [less info](#)

[edit](#) [privileges](#) [preferences](#) [flair](#) [apps](#) [my logins](#) | [meta user](#) | [network profile](#)



217,632
reputation

● 36 ● 295 ● 457

bio

website [karwin.com](#)

location California

age 47

visits

member for 6 years

visited 1922 days, 393 consecutive

seen 9 secs ago

stats

profile views 26,882

helpful flags 20

private

email [bill@karwin.com](#)

real name Bill Karwin

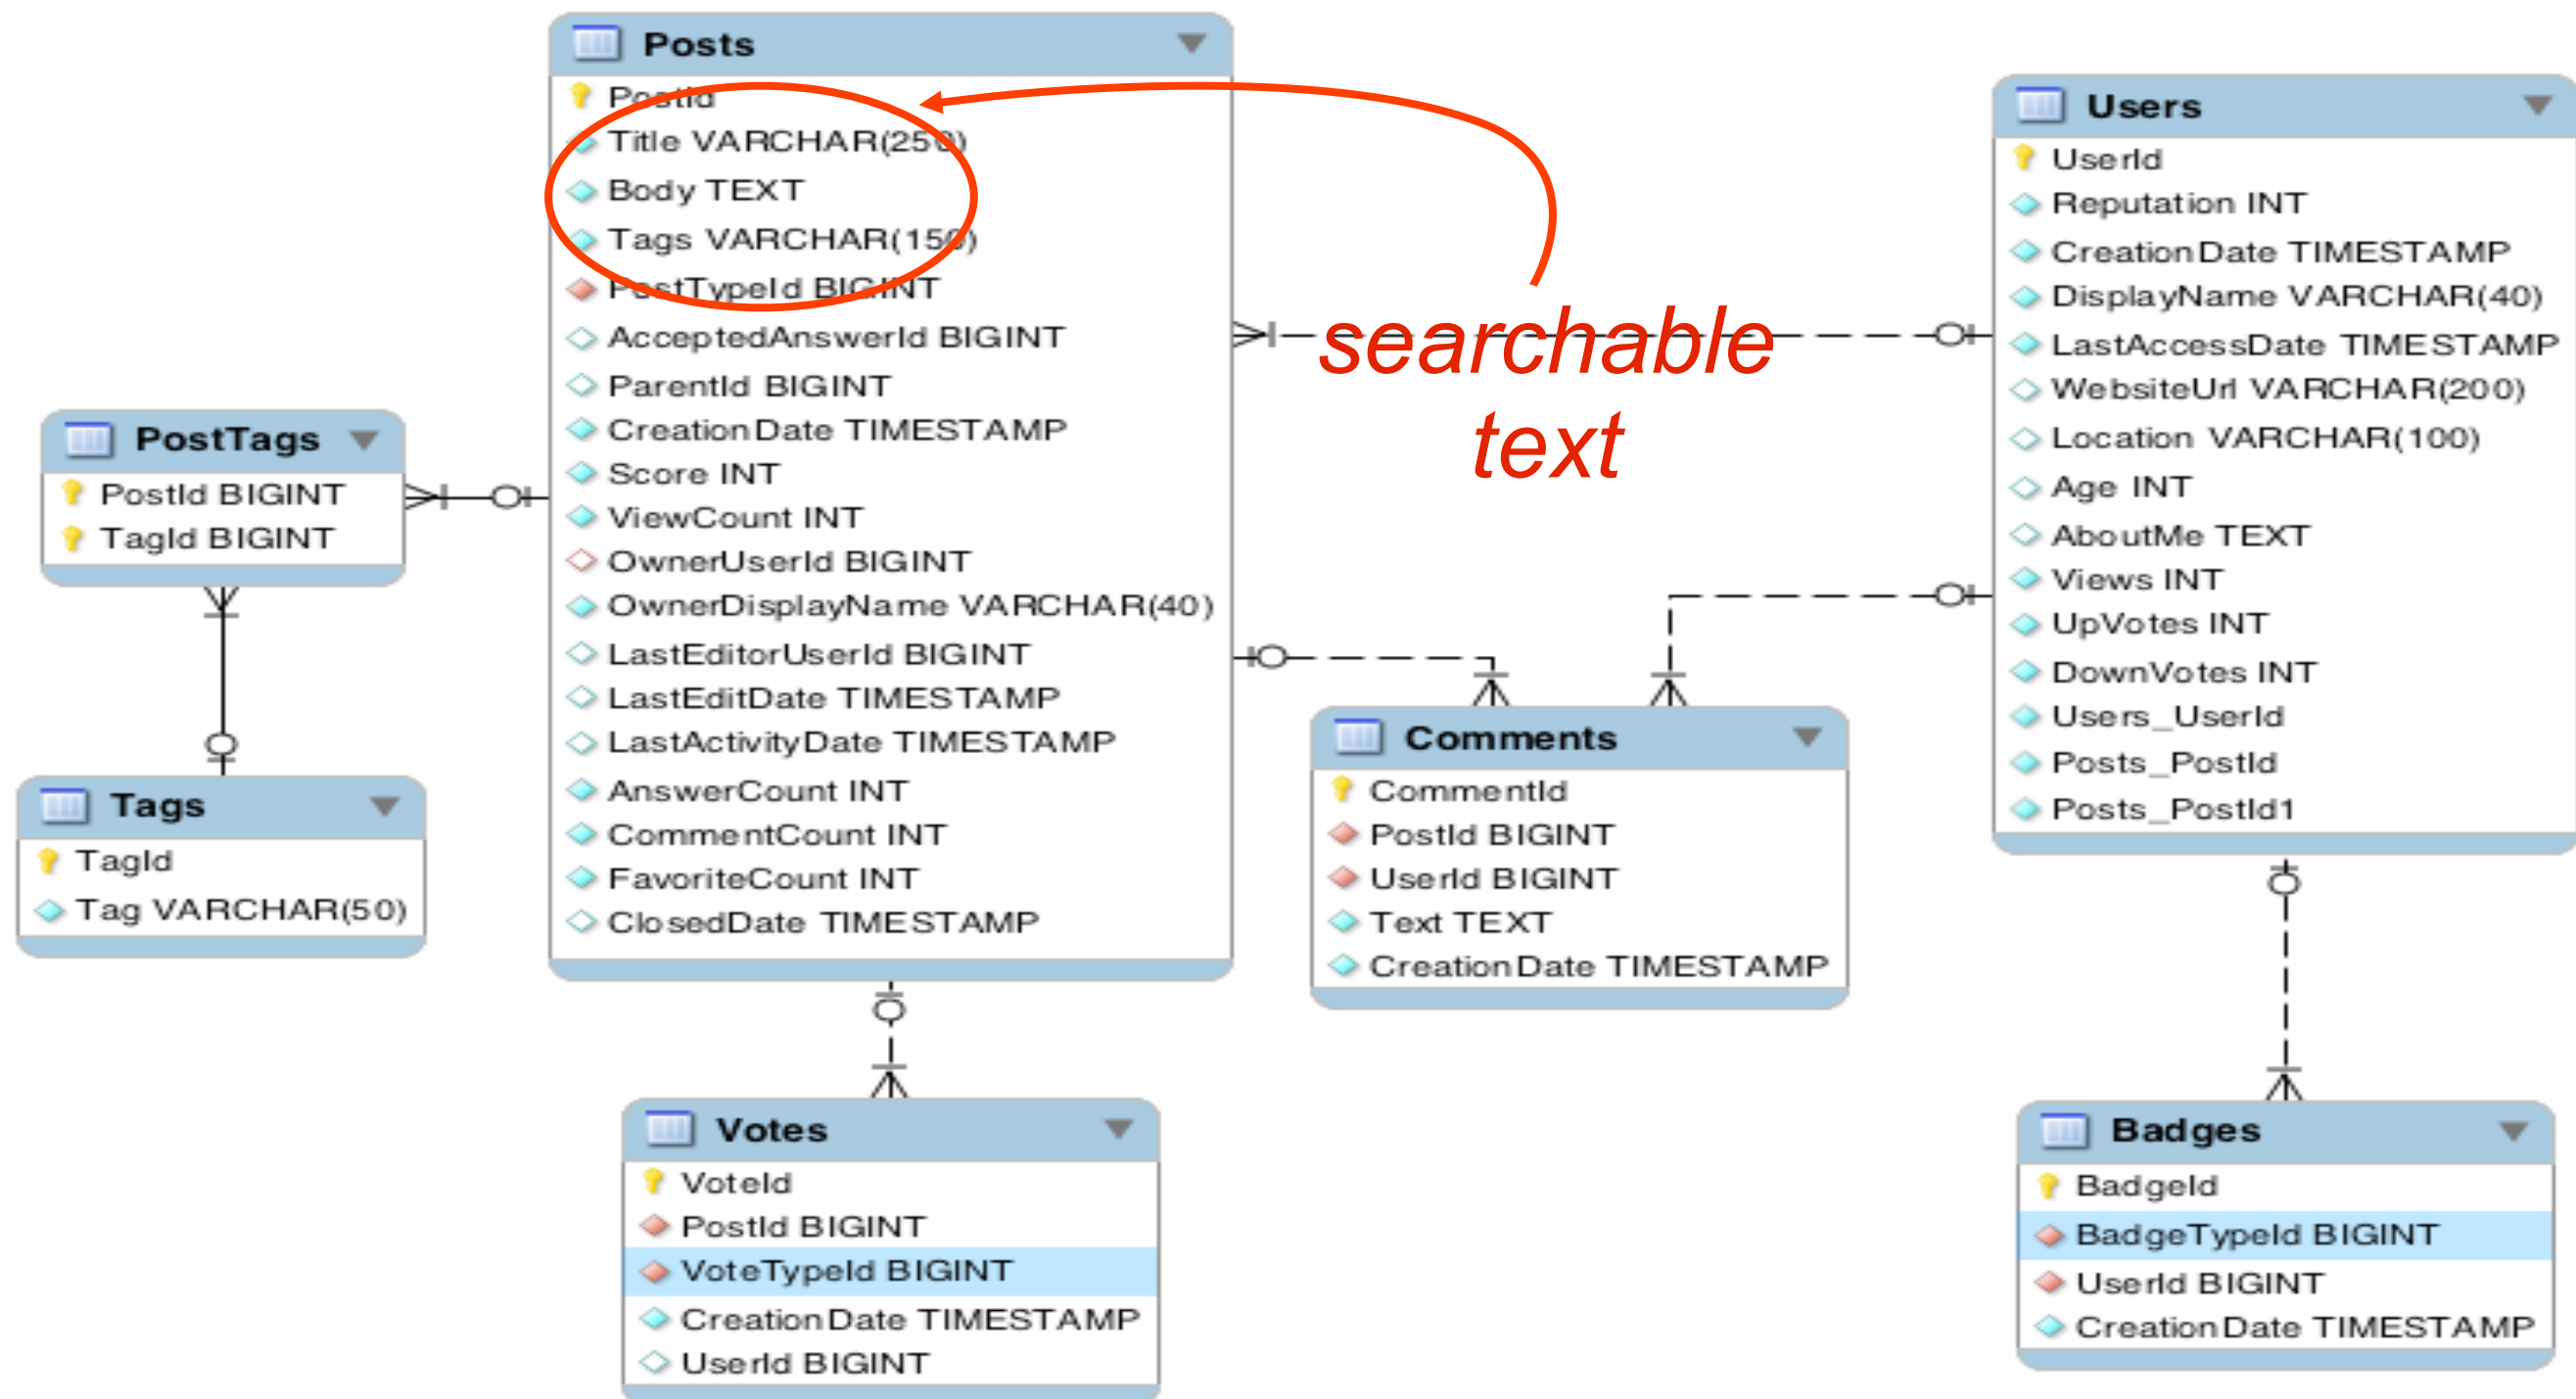
I'm Senior Knowledge Manager at Percona, a services company specializing in the MySQL database.

I've been a software engineer since 1987, and my specialty is as an SQL maven. I also have experience programming in Java, PHP, Perl, C, JavaScript, and I have many other coding skills.

I've written a book, [SQL Antipatterns: Avoiding the Pitfalls of Database Programming](#) from Pragmatic Bookshelf, based on the most common SQL problems I've answered on Stack Overflow and other forums, mailing lists, and newsgroups over the past 15 years.

I am an [Oracle ACE](#).

StackOverflow ER diagram



The Baseline: Naive Search Predicates

Some people, when confronted with a problem, think
“I know, I’ll use regular expressions.”
Now they have two problems.

— Jamie Zawinsky

Accuracy issue

- Irrelevant or false matching words 'one', 'money', 'prone', etc.:

```
SELECT * FROM Posts  
WHERE Body LIKE '%one%'
```

- Regular expressions in MySQL support escapes for word boundaries:

```
SELECT * FROM Posts  
WHERE Body RLIKE '[[<:]]one[[:>:]]'
```

Performance issue

- LIKE with wildcards:

```
SELECT * FROM Posts
WHERE title LIKE '%performance%'
      OR body LIKE '%performance%'
      OR tags LIKE '%performance%';
```


97 sec



- POSIX regular expressions:

```
SELECT * FROM Posts
WHERE title RLIKE '[[<:]]performance[[:>:]]'
      OR body RLIKE '[[<:]]performance[[:>:]]'
      OR tags RLIKE '[[<:]]performance[[:>:]]';
```

655 sec



Why so slow?

```
CREATE TABLE TelephoneBook (  
    FullName VARCHAR(50));
```

```
CREATE INDEX name_idx ON TelephoneBook  
    (FullName);
```

```
INSERT INTO TelephoneBook VALUES  
    ( 'Riddle, Thomas' ),  
    ( 'Thomas, Dean' );
```

Why so slow?

- Search for all with last name “Thomas”

```
SELECT * FROM telephone_book  
WHERE full_name LIKE 'Thomas%'
```

uses index



- Search for all with first name “Thomas”

```
SELECT * FROM telephone_book  
WHERE full_name LIKE '%Thomas'
```

can't use index



Because:



*B-Tree indexes can't
search for substrings*

FULLTEXT in MyISAM

FULLTEXT in InnoDB

Apache Solr

Sphinx Search

Trigraphs

FULLTEXT in MyISAM

FULLTEXT Index with MyISAM

- Special index type for MyISAM
- Integrated with SQL queries
- Indexes always in sync with data
- Balances features vs. speed vs. space
- Testing: MySQL Community Edition 5.7.5

Build Index on Data (MyISAM)

```
mysql> CREATE FULLTEXT INDEX PostText  
ON Posts(title, body, tags);
```



time: 30 min, 10 sec

Size of Index (MyISAM)

```
mysql>SHOW TABLE STATUS LIKE 'Posts'\G
```

Name: posts

Engine: MyISAM

Rows: 8000000


Avg_row_length: 927

Data_length: 7417899480 (6.91GB)

Index_length: 2803019776 (Δ = 2.50GB)

Querying

```
SELECT * FROM Posts  
WHERE MATCH(column(s))  
AGAINST('query pattern');
```



*must include all columns
of your index, in the
order you defined*

Natural Language Mode (MyISAM)

- Searches concepts with free text queries:

```
SELECT * FROM Posts  
WHERE MATCH(title, body, tags )  
AGAINST('mysql performance'  
IN NATURAL LANGUAGE MODE)  
LIMIT 100;
```



*time with index:
183 milliseconds*

Query Profile: Natural Language Mode (MyISAM)

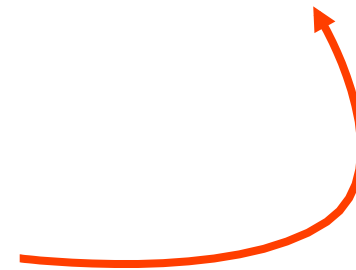
Status	Duration
starting	0.000112
checking permissions	0.000012
Opening tables	0.000155
checking permissions	0.000004
checking permissions	0.000032
init	0.000007
checking permissions	0.000070
System lock	0.000012
optimizing	0.000009
statistics	0.000028
preparing	0.000009
FULLTEXT initialization	0.180492
executing	0.000012
Sending data	0.002302
end	0.000004
query end	0.000005
closing tables	0.000017
freeing items	0.000432
cleaning up	0.000090

Boolean Mode (MyISAM)

- Searches words using mini-language:

```
SELECT * FROM Posts  
WHERE MATCH(title, body, tags)  
AGAINST(' +mysql +performance'  
IN BOOLEAN MODE)  
LIMIT 100;
```

*time with index:
10 milliseconds*



Query Profile: Boolean Mode (MyISAM)

Status	Duration
starting	0.000123
checking permissions	0.000009
Opening tables	0.000153
checking permissions	0.000005
checking permissions	0.000033
init	0.000007
checking permissions	0.000070
System lock	0.000012
optimizing	0.000009
statistics	0.000027
preparing	0.000009
FULLTEXT initialization	0.000025
executing	0.000008
Sending data	0.009507
end	0.000009
query end	0.000008
closing tables	0.000017
freeing items	0.000107
cleaning up	0.000024

FULLTEXT in InnoDB

FULLTEXT Index with InnoDB

- Usage very similar to FULLTEXT in MyISAM
- Integrated with SQL queries
- Indexes always* in sync with data
- Read the blogs for more details:
 - <http://blogs.innodb.com/wp/2011/07/overview-and-getting-started-with-innodb-fts/>
 - <http://blogs.innodb.com/wp/2011/07/innodb-full-text-search-tutorial/>
 - <http://blogs.innodb.com/wp/2011/07/innodb-fts-performance/>
 - <http://blogs.innodb.com/wp/2011/07/difference-between-innodb-fts-and-mysam-fts/>
- Testing: MySQL Community Edition 5.7.5

Build Index on Data (InnoDB)

- Relatively new code; you might see problems:

```
mysql> CREATE FULLTEXT INDEX PostText  
ON Posts(title, body, tags);
```

```
ERROR 2013 (HY000): Lost connection to  
MySQL server during query
```

Build Index on Data (InnoDB)

- Solution: define a primary key column called ``FTS_DOC_ID`` explicitly:

```
mysql> ALTER TABLE Posts  
CHANGE COLUMN PostId  
`FTS_DOC_ID` BIGINT UNSIGNED;
```

```
mysql> CREATE FULLTEXT INDEX PostText  
ON Posts(title, body, tags);
```



time: 30 min, 19 sec

Size of Index (InnoDB)

```
mysql>SHOW TABLE STATUS LIKE 'Posts'\G
```

Name: posts

Engine: InnoDB

Rows: 6877702

Avg_row_length: 1427

Data_length: 9817817088 (9.14GB)

Index_length: 0 *



2.62 GB on disk

Natural Language Mode (InnoDB)

- Searches concepts with free text queries:

```
SELECT * FROM Posts  
WHERE MATCH(title, body, tags)  
AGAINST('mysql performance'  
IN NATURAL LANGUAGE MODE)  
LIMIT 100;
```



*time with index:
610 milliseconds*

Query Profile: Natural Language Mode (InnoDB)

Status	Duration
starting	0.000185
checking permissions	0.000009
Opening tables	0.000154
checking permissions	0.000004
checking permissions	0.000033
init	0.000010
checking permissions	0.000072
System lock	0.000010
optimizing	0.000011
statistics	0.000027
preparing	0.000010
FULLTEXT initialization	0.404054
executing	0.000013
Sending data	0.143711
end	0.000014
query end	0.000011
closing tables	0.000011
freeing items	0.063529
cleaning up	0.000059

Boolean Mode (InnoDB)

- Searches words using mini-language:

```
SELECT * FROM Posts  
WHERE MATCH(title, body, tags)  
AGAINST(' +mysql +performance'  
IN BOOLEAN MODE)  
LIMIT 100;
```

*time with index:
323 milliseconds*



Query Profile: Boolean Mode (InnoDB)

Status	Duration
starting	0.000128
checking permissions	0.000008
Opening tables	0.000155
checking permissions	0.000004
checking permissions	0.000033
init	0.000007
checking permissions	0.000072
System lock	0.000010
optimizing	0.000009
statistics	0.000027
preparing	0.000010
FULLTEXT initialization	0.313276
executing	0.000010
Sending data	0.008098
end	0.000008
query end	0.000011
closing tables	0.000012
freeing items	0.001478
cleaning up	0.000024

Apache Solr

Apache Solr

- <http://lucene.apache.org/solr/>
- Formerly known as Lucene, started 2001
- Apache License
- Java implementation
- Web service architecture
- Many sophisticated search feature
- Testing: Apache Solr 4.10.1, Java 8

DataImportHandler

- *conf/solrconfig.xml:*

```
. . .  
<requestHandler name="/dataimport"  
  class="solr.DataImportHandler">  
  <lst name="defaults">  
    <str name="config">solr-data-config.xml</str>  
  </lst>  
</requestHandler>  
. . .
```

DataImportHandler

- *conf/data-config.xml*:

```
<dataConfig>
  <dataSource type="JdbcDataSource"
    driver="com.mysql.jdbc.Driver"
    url="jdbc:mysql://localhost/testpattern?useUnicode=true"
    batchSize="-1"
    user="xxxx"
    password="xxxx" />
  <document>
    <entity name="id"
      query="SELECT PostId, ParentId, Title, Body, Tags FROM Posts">
    </entity>
  </document>
</dataConfig>
```

*extremely important
to avoid buffering the
whole query result!*

DataImportHandler

- *conf/schema.xml:*

```
. . .
<fields>
  <field name="Id" type="string" indexed="true" stored="true" required="true" />
  <field name="ParentId" type="string" indexed="true" stored="true" required="false" />
  <field name="Title" type="text_general" indexed="false" stored="false"
    required="false" />
  <field name="Body" type="text_general" indexed="false" stored="false" required="false"
    >
  <field name="Tags" type="text_general" indexed="false" stored="false" required="false"
    >

  <field name="text" type="text_general" indexed="true" stored="false" multiValued="true"
    >
</fields>

<uniqueKey>PostId</uniqueKey>
<defaultSearchField>text</defaultSearchField>

<copyField source="Title" dest="text"/>
<copyField source="Body" dest="text"/>
<copyField source="Tags" dest="text"/>
. . .
```

Insert Data into Index (Solr)

The screenshot displays the Apache Solr Admin interface in a web browser. The browser's address bar shows the URL `localhost:8983/solr/#/stackoverflow/dataimport//dataimport`. On the left sidebar, the 'Dataimport' menu item is selected. The main content area shows the status of a data import operation. A green banner at the top indicates 'Indexing completed. Added/Updated: 8,000,000 documents. Deleted 0 documents. (Duration: 18m 52s)'. Below this, a table shows the progress: 'Requests: 1 (0/s), Fetched: 8,000,000 (7,067/s), Skipped: 0, Processed: 8,000,000 (7,067/s)'. The 'Duration: 18m 52s' is circled in red. The left sidebar contains various navigation options: Dashboard, Logging, Core Admin, Java Properties, Thread Dump, stackoverflow (selected), Overview, Analysis, Dataimport (selected), Documents, Files, Ping, Plugins / Stats, Query, Replication, and Schema Browser. The bottom of the page features links to Documentation, Issue Tracker, IRC Channel, Community forum, and Solr Query Syntax.

Apache Solr Admin

localhost:8983/solr/#/stackoverflow/dataimport//dataimport

Apache Solr

Dashboard

Logging

Core Admin

Java Properties

Thread Dump

stackoverflow

Overview

Analysis

Dataimport

Documents

Files

Ping

Plugins / Stats

Query

Replication

Schema Browser

/dataimport

Command: full-import

Verbose: ☐

Clean: ☒

Commit: ☒

Optimize: ☐

Debug: ☐

Entity:

Start, Rows: 0

Custom Parameters: key1=val1&key2=val2

Execute

Refresh Status

Auto-Refresh Status: ☒

Last Update: 17:31:33

Indexing completed. Added/Updated: 8,000,000 documents. Deleted 0 documents. (Duration: 18m 52s)

Requests: 1 (0/s), Fetched: 8,000,000 (7,067/s), Skipped: 0, Processed: 8,000,000 (7,067/s)

Started: less than a minute ago

Raw Status-Output

Configuration

Debug-Mode

Reload

Documentation

Issue Tracker

IRC Channel

Community forum

Solr Query Syntax

time: 18 min 52 sec

Size of Index (Solr)

The screenshot displays the Apache Solr Admin interface for a cluster named 'stackoverflow'. The left sidebar contains navigation links: Dashboard, Logging, Core Admin, Java Properties, Thread Dump, Overview (selected), Analysis, Dataimport, Documents, Files, Ping, Plugins / Stats, Query, Replication, and Schema Browser. The main content area is divided into several sections:

- Statistics:** Last Modified: 5 minutes ago, Num Docs: 8000000, Max Doc: 8000000, Heap Memory: 17424896, Usage: Deleted Docs: 0, Version: 1114, Segment Count: 36, Optimized: , Current: .
- Instance:** CWD: c:\Solr\example, Instance: c:\Solr\example\example-DIH\solr\stackoverflow, Data: C:\Solr\example\example-DIH\solr\stackoverflow\data, Index: C:\Solr\example\example-DIH\solr\stackoverflow\data\index, Impl: org.apache.solr.core.NRTCachingDirectoryFactory.
- Replication (Master):** A table showing replication details for the master node.
- Healthcheck:** Ping request handler is not configured with a healthcheck file.
- Admin Extra:** A section for additional administrative actions.

The 'Replication (Master)' table is highlighted with a red circle, and an arrow points from the text 'size: 2.49 GB' to the 'Size' column of the 'Master (Searching)' row.

	Version	Gen	Size
Master (Searching)	1413937513151	76	2.49 GB
Master (Replicable)	1413937513151	76	-

size: 2.49 GB

Searching Solr

- <http://localhost:8983/solr/stackoverflow/query?q=mysql+AND+performance>

Request-Handler (qt)

/select

common

q

mysql AND performance

fq

sort

start, rows

0 10

fl

df

Raw Query Parameters

key1=val1&key2=val2

wt

json

☒ indent

☐ debugQuery

☐ dismax

☐ edismax

☐ hl

☐ facet

☐ spatial

☐ spellcheck

Execute Query

```
{
  "responseHeader": {
    "status": 0,
    "QTime": 0,
    "params": {
      "q": "mysql AND performance",
      "indent": "true",
      "wt": "json",
      "_": "1413938439009"
    }
  },
  "response": {
    "numFound": 7951,
    "start": 0,
    "docs": [
      {
        "id": "761733",
        "_version_": 1482619676739502000
      },
      {
        "id": "2481528",
        "_version_": 1482619882841309200
      },
      {
        "id": "4184107",
        "_version_": 1482620128668418000
      },
      {
        "id": "3820995",
        "_version_": 1482620084826407000
      },
      {
        "id": "3725227",
        "_version_": 1482620073021538600
      },
      {
        "id": "5594157",
        "_version_": 1482620358333825000
      }
    ]
  }
}
```

time: 60-1700ms

Query results are cached (like MySQL Query Cache), so a given search returns much faster (< 1ms) on subsequent execution

Sphinx Search

Sphinx Search

- <http://sphinxsearch.com/>
- Started in 2001
- GPLv2 license
- C++ implementation
- SphinxSE storage engine for MySQL
- Supports MySQL protocol, SQL-like queries
- Many sophisticated search features
- Testing: Sphinx Search 2.2.5

sphinx.conf

```
source src1
{
  type = mysql
  sql_host = localhost
  sql_user = xxxx
  sql_pass = xxxx
  sql_db = testpattern
  sql_query = SELECT PostId, ParentId, Title,
    Body, Tags FROM Posts
  sql_query_info = SELECT * FROM Posts \
    WHERE PostId=$id
}
```

sphinx.conf

```
index test1
{
  source = src1
  path = C:\Sphinx\data
}
```

Insert Data into Index (Sphinx)

```
Administrator: C:\Windows\system32\cmd.exe

c:\Sphinx\bin>indexer.exe -c ../sphinx.conf --verbose test1
Sphinx 2.2.5-id64-release (r4825)
Copyright (c) 2001-2014, Andrew Aksyonoff
Copyright (c) 2008-2014, Sphinx Technologies Inc (http://sphinxsearch.com)

using config file ' ../sphinx.conf'...
indexing index 'test1'...
WARNING: attribute 'group_id' not found - IGNORING
WARNING: attribute 'date_added' not found - IGNORING
WARNING: Attribute count is 0: switching to none docinfo
collected 8000000 docs, 6827.1 MB
sorted 1134.6 Mhits, 100.0% done
total 8000000 docs, 6827097596 bytes
total 1071.670 sec, 6370519 bytes/sec, 7464.98 docs/sec
total 8941 reads, 280.311 sec, 495.0 kb/call avg, 31.3 msec/call avg
total 9111 writes, 5.472 sec, 932.8 kb/call avg, 0.6 msec/call avg

c:\Sphinx\bin>
```

time: 17 min 52 sec

Index Size (Sphinx)

Name	Date modified	Type	Size
rt.kill	10/21/2014 6:54 PM	KILL File	1 KB
rt.lock	10/21/2014 6:54 PM	LOCK File	0 KB
rt.meta	10/21/2014 6:54 PM	META File	1 KB
rt.ram	10/21/2014 6:54 PM	RAM File	1 KB
test1.spa	10/21/2014 6:28 PM	SPA File	0 KB
test1.spd	10/21/2014 6:39 PM	SPD File	2,335,523 KB
test1.spe	10/21/2014 6:39 PM	SPE File	19,928 KB
test1.sph	10/21/2014 6:39 PM	SPH File	1 KB
test1.spi	10/21/2014 6:39 PM	SPI File	84,380 KB
test1.spk	10/21/2014 6:28 PM	SPK File	0 KB
test1.spl	10/21/2014 6:54 PM	Shockwave Flash ...	0 KB
test1.spm	10/21/2014 6:28 PM	SPM File	0 KB
test1.spp	10/21/2014 6:39 PM	SPP File	1,632,791 KB
test1.sps	10/21/2014 6:39 PM	SPS File	1 KB
test1.tmps	10/21/2014 6:39 PM	TMPS File	1 KB

15 items selected Date modified: 10/21/2014 6:54 PM Date created: 10/21/2014 6:54 PM
Size: 3.88 GB

size: 3.88 GB

Querying index

```
$ mysql --port 9306
```

```
Server version: 2.2.5-id64-release (r4825)
```

```
mysql> SELECT * FROM test1 WHERE MATCH('mysql performance');
```

```
+-----+  
| id      |  
+-----+  
| 6016856 |  
| 4207641 |  
| 2656325 |  
| 7192928 |  
| 8118235 |  
.  
.  
.
```

```
20 rows in set (0.04 sec)
```

Querying index

```
mysql> SHOW META;
```

Variable_name	Value
total	1000
total_found	8340
time	0.037
keyword[0]	mysql
docs[0]	179579
hits[0]	404247
keyword[1]	performance
docs[1]	158433
hits[1]	227427

time: 37ms

Trigraphs

Trigraphs Overview

- Not very fast, but still better than LIKE / RLIKE
- Generic, portable SQL solution
- No dependency on version, storage engine, third-party technology

Three-Letter Sequences

```
CREATE TABLE AtoZ (  
  c      CHAR(1),  
  PRIMARY KEY (c));  
  
INSERT INTO AtoZ (c)  
VALUES ('a'), ('b'), ('c'), ...  
  
CREATE TABLE Trigraphs (  
  Tri     CHAR(3),  
  PRIMARY KEY (Tri));  
  
INSERT INTO Trigraphs (Tri)  
SELECT CONCAT(t1.c, t2.c, t3.c)  
FROM AtoZ t1 JOIN AtoZ t2 JOIN AtoZ t3;
```

Insert Data Into Index

```
my $sth = $dbh1->prepare("SELECT * FROM Posts") or die $dbh1->errstr;
$sth->execute() or die $dbh1->errstr;
$dbh2->begin_work;
my $i = 0;
while (my $row = $sth->fetchrow_hashref ) {
    my $text = lc(join('|', ($row->{title}, $row->{body}, $row->{tags})));
    my %tri;
    map($tri{$_}=1, ( $text =~ m/[[[:alpha:]]{3}/g ));
    next unless %tri;
    my $tuple_list = join(", ", map("('$_', $row->{postid})", keys %tri));
    my $sql = "INSERT IGNORE INTO PostsTrigraph (tri, PostId) VALUES
    $tuple_list";
    $dbh2->do($sql) or die "SQL = $sql, ".$dbh2->errstr;
    if (++$i % 1000 == 0) {
        print ".";
        $dbh2->commit;
        $dbh2->begin_work;
    }
}
print ".\n";
$dbh2->commit;
```

*takes hours, and
creates a very large
number of rows*

Indexed Lookups

```
SELECT p.*  
FROM Posts p  
JOIN PostsTrigraph t1 ON  
  t1.PostId = p.PostId AND t1.Tri = 'mys'
```

Search Among Fewer Matches

```
SELECT p.*  
FROM Posts p  
JOIN PostsTrigraph t1 ON  
  t1.PostId = p.PostId AND t1.Tri = 'mys'  
JOIN PostsTrigraph t2 ON  
  t2.PostId = p.PostId AND t2.Tri = 'per'
```


Search Among Fewer Matches

```
SELECT p.*  
FROM Posts p  
JOIN PostsTrigraph t1 ON  
  t1.PostId = p.PostId AND t1.Tri = 'mys'  
JOIN PostsTrigraph t2 ON  
  t2.PostId = p.PostId AND t2.Tri = 'per'  
JOIN PostsTrigraph t3 ON  
  t3.PostId = p.PostId AND t3.Tri = 'for'
```

Search Among Fewer Matches

```
SELECT p.*
FROM Posts p
JOIN PostsTrigraph t1 ON
  t1.PostId = p.PostId AND t1.Tri = 'mys'
JOIN PostsTrigraph t2 ON
  t2.PostId = p.PostId AND t2.Tri = 'per'
JOIN PostsTrigraph t3 ON
  t3.PostId = p.PostId AND t3.Tri = 'for'
JOIN PostsTrigraph t4 ON
  t4.PostId = p.PostId AND t4.Tri = 'man'
```

Narrow Down Further

```
SELECT p.*
FROM Posts p
JOIN PostsTrigraph t1 ON
  t1.PostId = p.PostId AND t1.Tri = 'mys'
JOIN PostsTrigraph t2 ON
  t2.PostId = p.PostId AND t2.Tri = 'per'
JOIN PostsTrigraph t3 ON
  t3.PostId = p.PostId AND t3.Tri = 'for'
JOIN PostsTrigraph t4 ON
  t4.PostId = p.PostId AND t4.Tri = 'man'
WHERE CONCAT(p.title,p.body,p.tags) LIKE '%mysql%'
  AND CONCAT(p.title,p.body,p.tags) LIKE '%performance%';
```

Not Recommended

- Best query performance was still > 15 sec.
- Specialized fulltext search technology is much better, so trigraphs are useful only when portable, standard SQL is the only solution allowed.



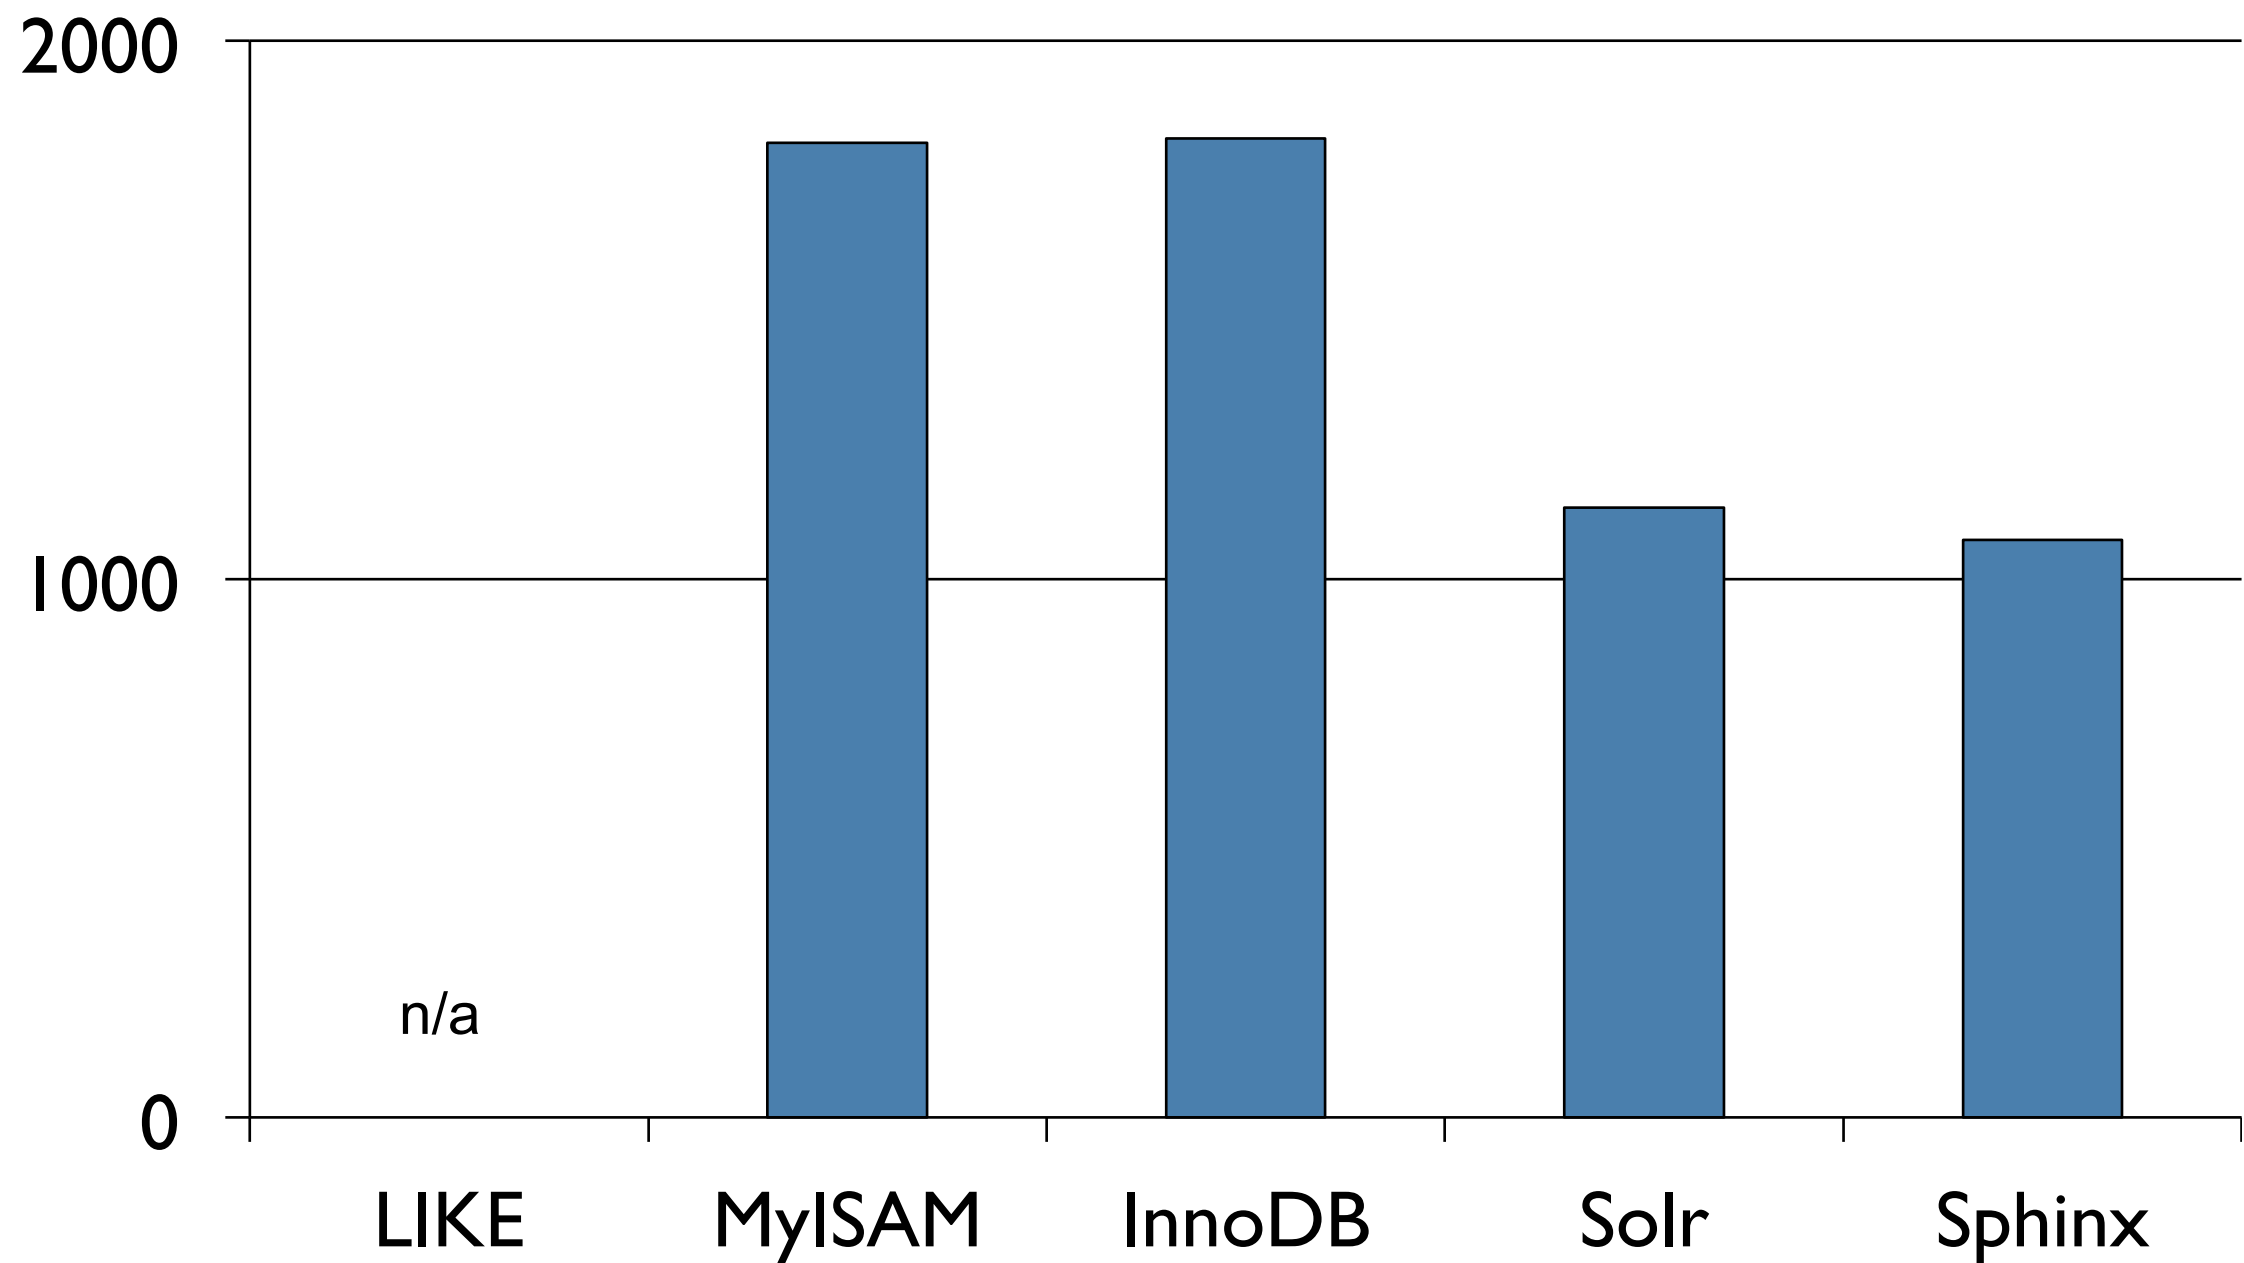
And the winner is...



Time to Build Index on Data

LIKE expression	n/a
FULLTEXT MyISAM	30 min, 10 sec
FULLTEXT InnoDB	30 min, 19 sec
Apache Solr	18 min, 52 sec
Sphinx Search	17 min, 52 sec

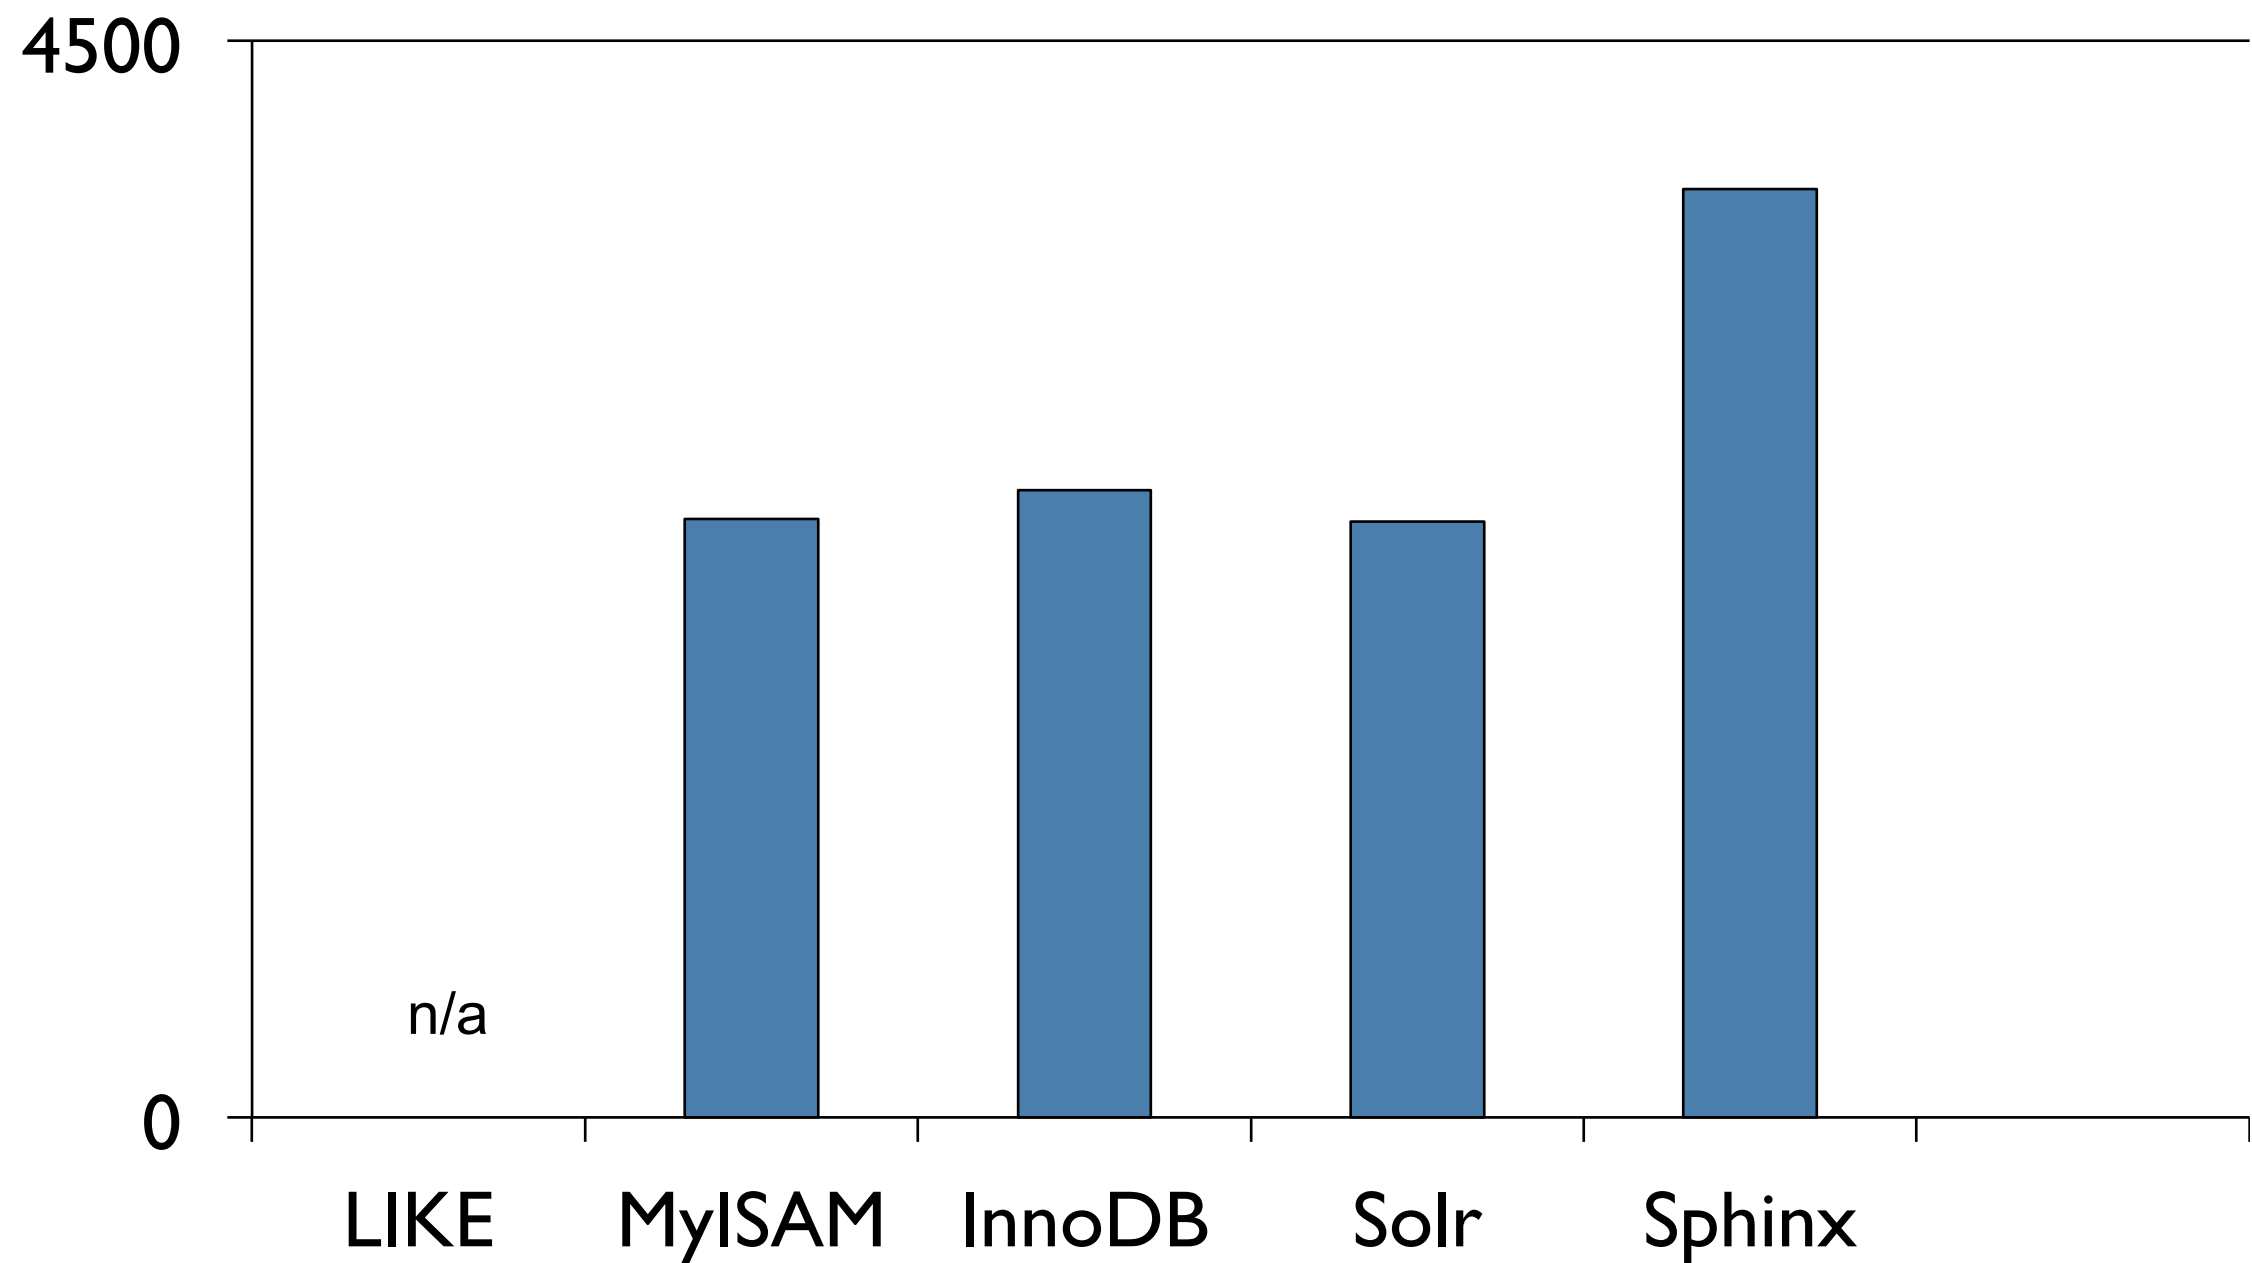
Build Index on Data (sec)



Index Storage

LIKE expression	n/a
FULLTEXT MyISAM	2.50 GB
FULLTEXT InnoDB	2.62 GB
Apache Solr	2.49 GB
Sphinx Search	3.88 GB

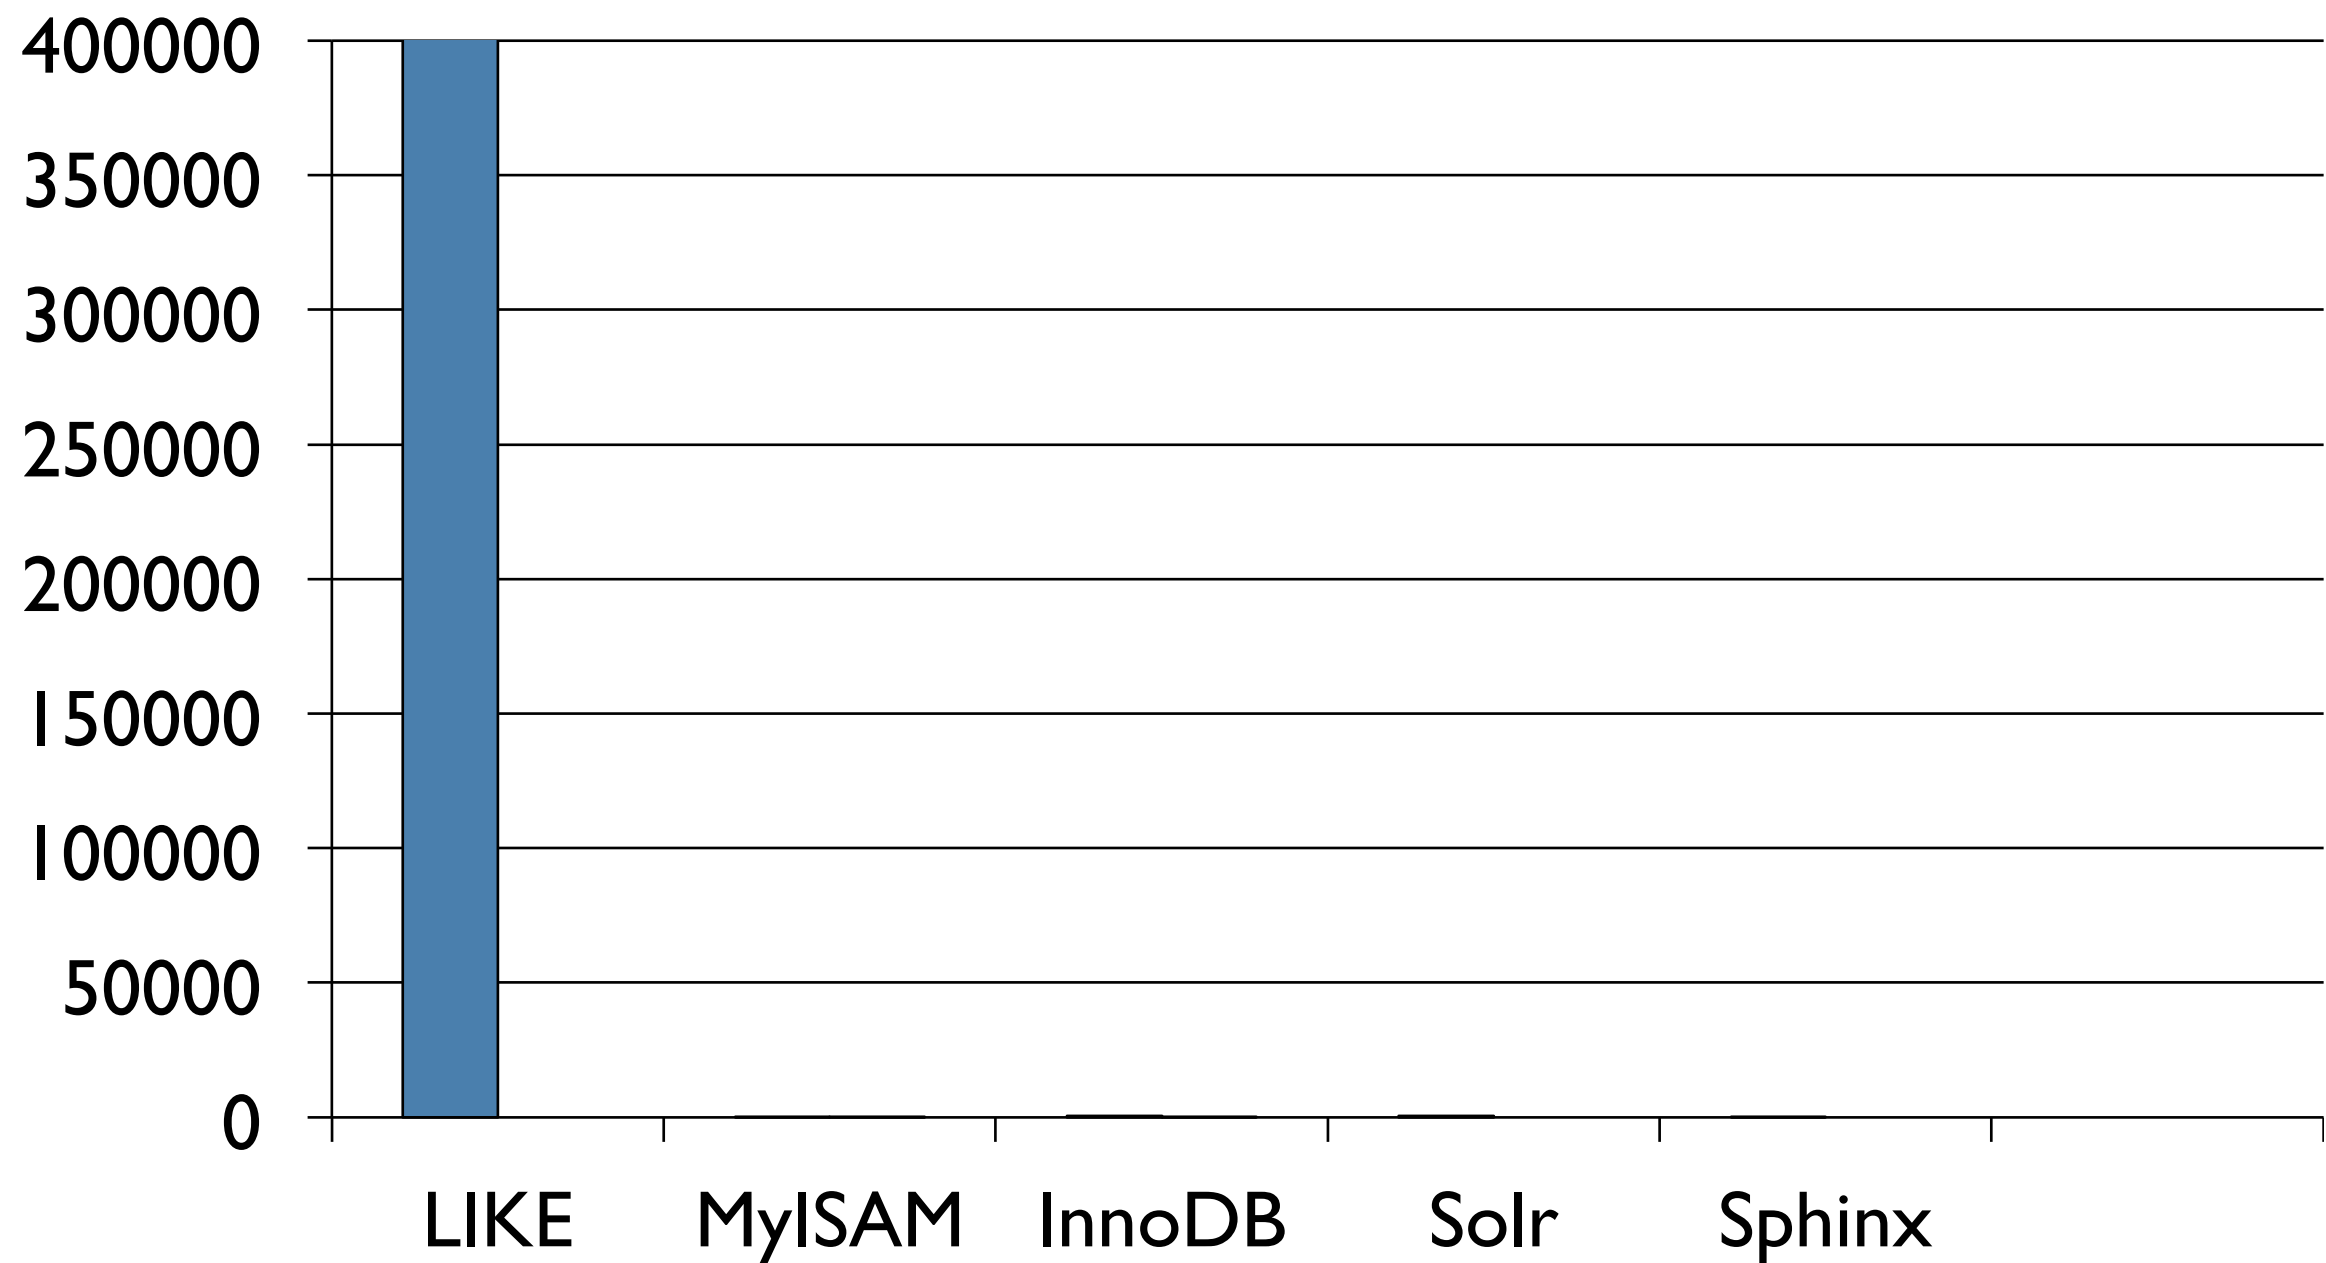
Index Storage (MiB)



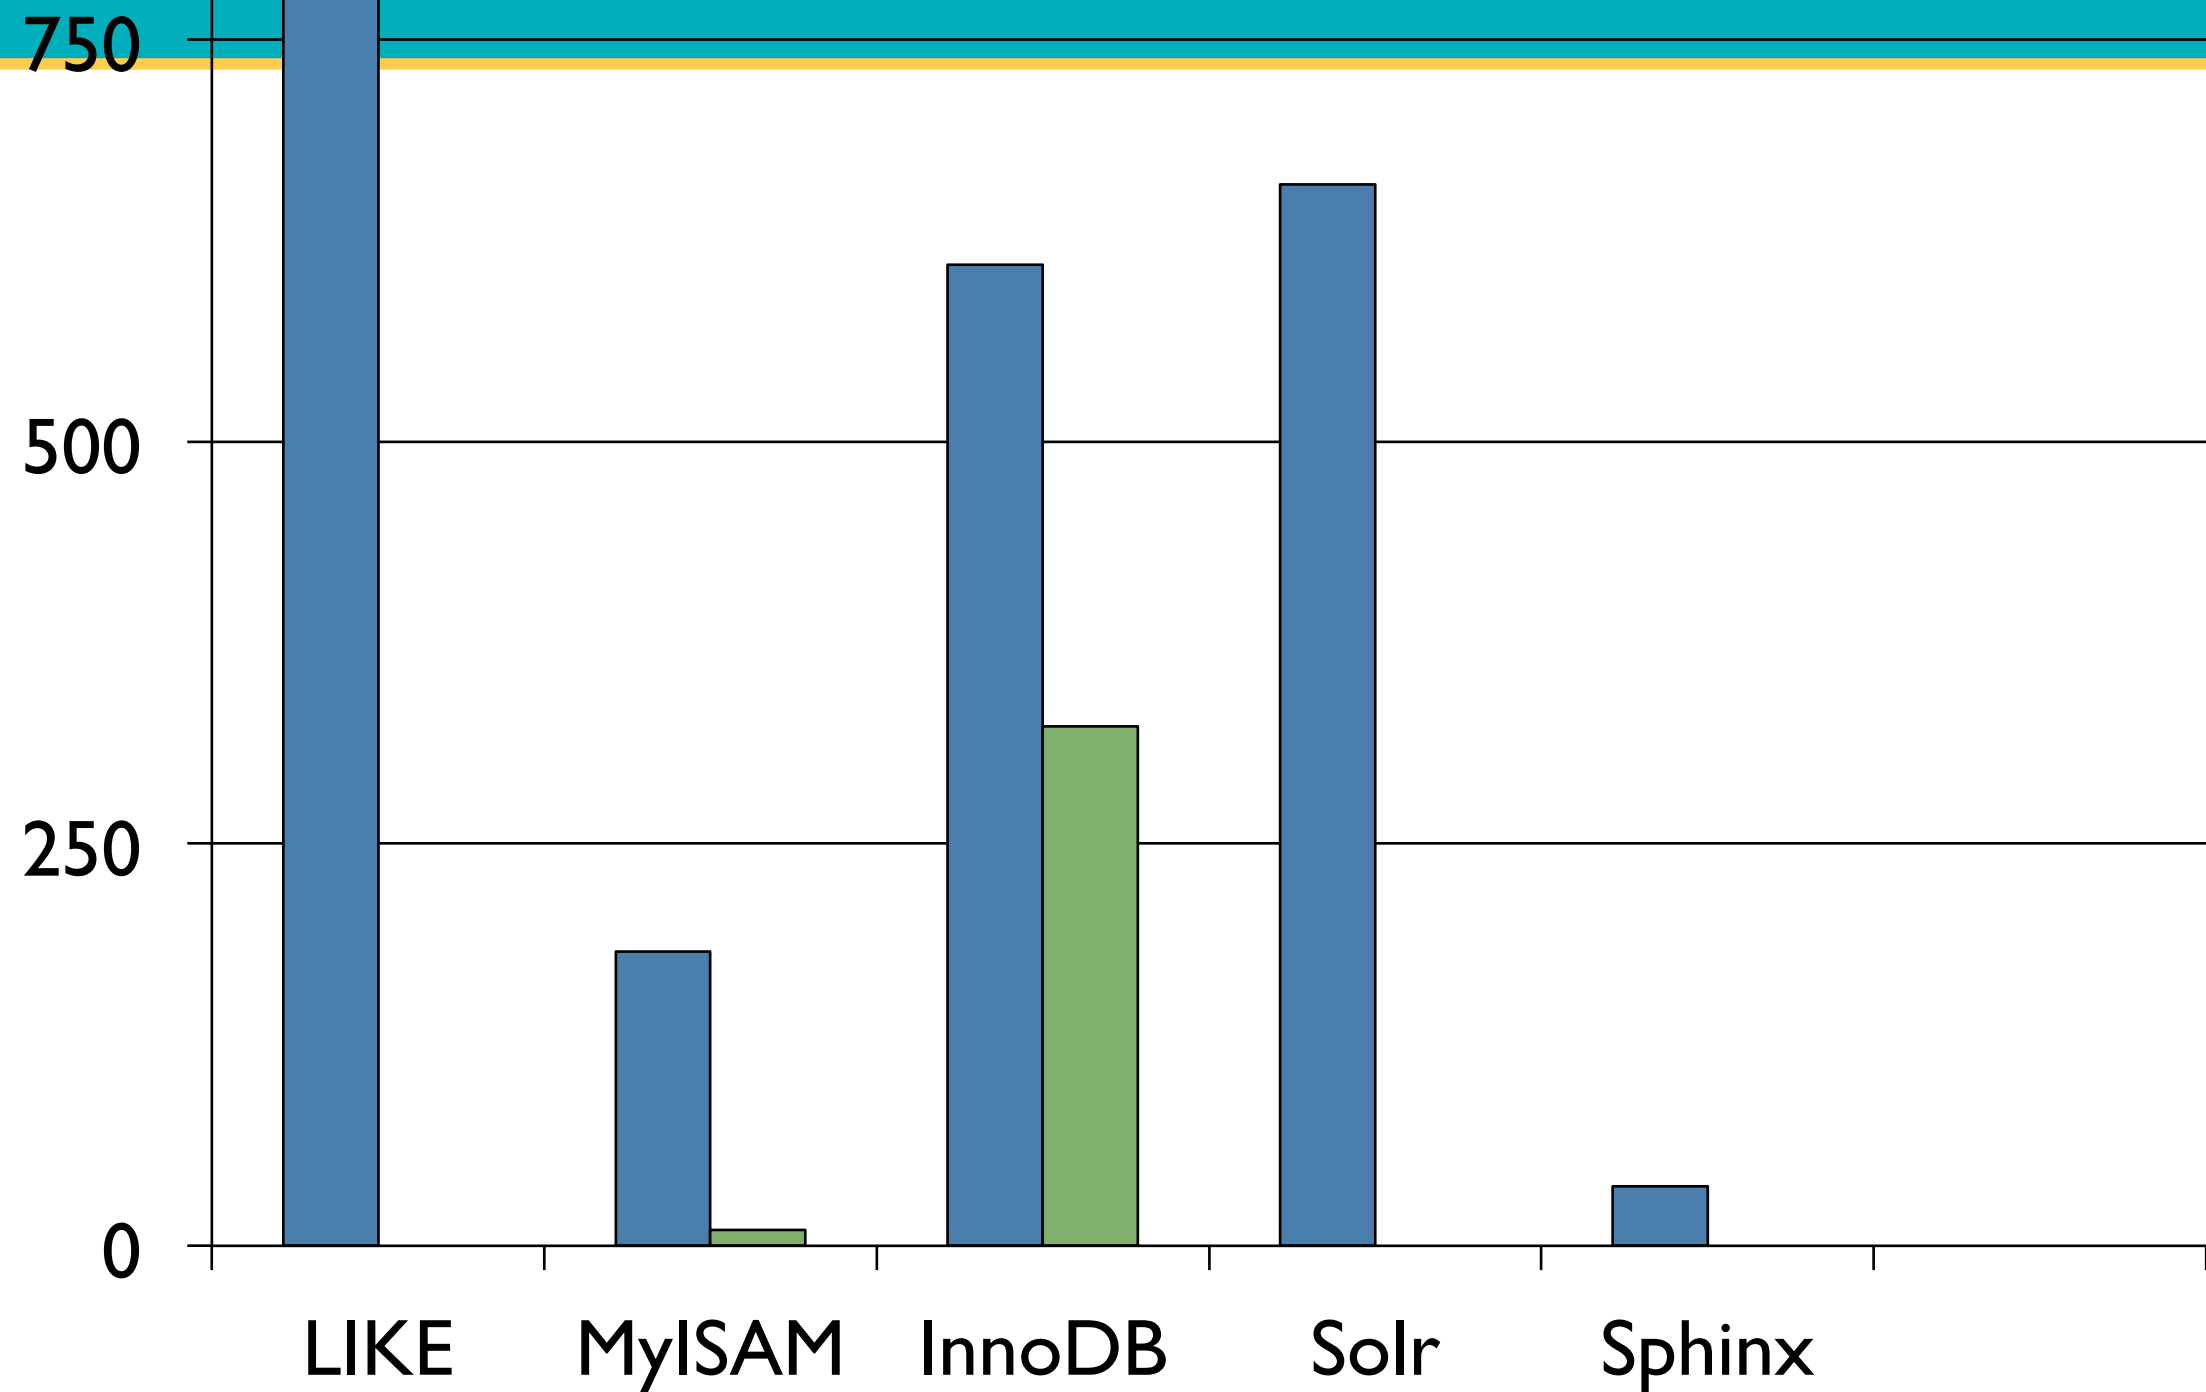
Query Speed

LIKE / RLIKE	97000-655000ms
FULLTEXT MyISAM	10-183ms
FULLTEXT InnoDB	323-610ms
Apache Solr	60-1700ms
Sphinx Search	37ms

Query Speed (ms)



Query Speed (ms)



Bottom Line

	<i>build</i>	<i>storage</i>	<i>query</i>	<i>solution</i>
LIKE expression	0	0	97,000ms	SQL
RLIKE expression	0	0	655,000ms	SQL
FULLTEXT MyISAM	30:10	2.50GB	183ms	MySQL
FULLTEXT InnoDB	30:19	2.62GB	323ms	MySQL 5.6
Apache Solr	18:52	2.49GB	660ms	Java
Sphinx Search	17:52	3.88GB	37ms	C++

Final Thoughts

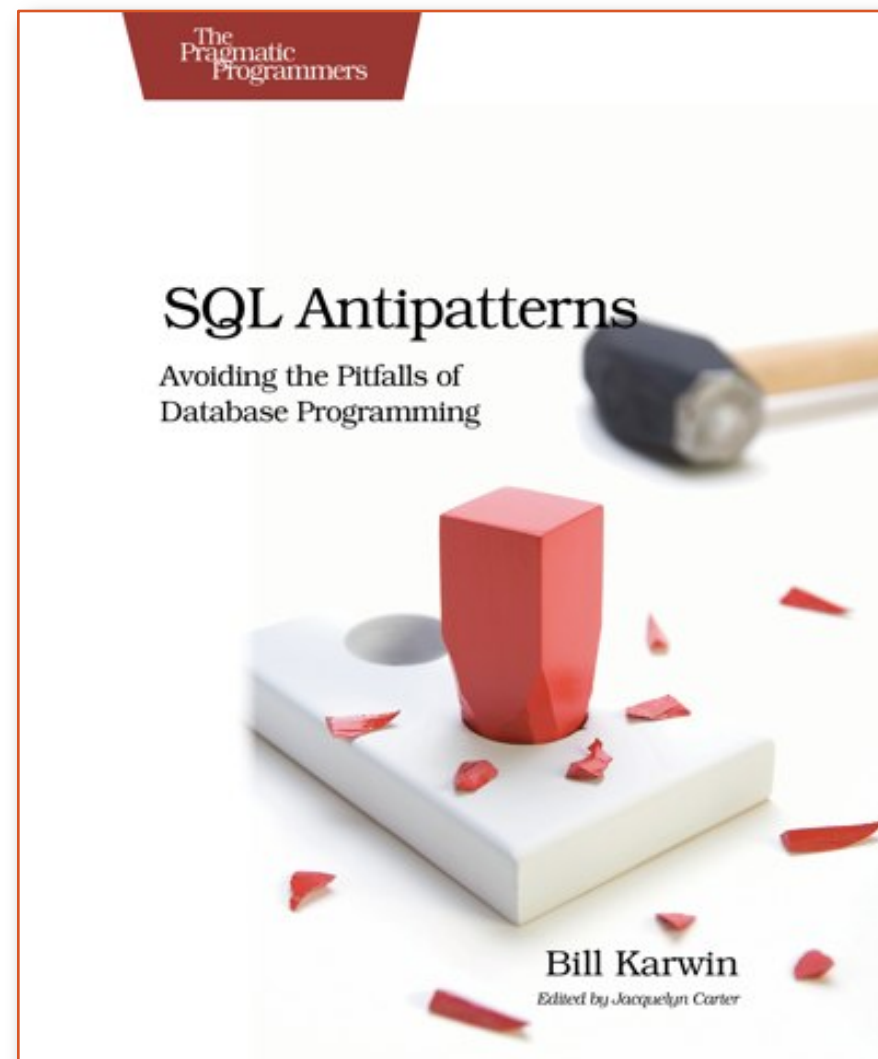
- Third-party search engines are complex to keep in sync with data, and adding another type of server adds more operations work for you.
- Built-in FULLTEXT indexes are therefore useful even if they are not absolutely the fastest.

Final Thoughts

- Different search implementations may return different results, so you should evaluate what works best for your project.

Final Thoughts

- *Any* indexed search solution is orders of magnitude better than LIKE!



<http://www.pragprog.com/titles/bksqla/>

London, November 3-4, 2014
Santa Clara, April 13-16, 2015

www.percona.com/live



Super-saver registration for Santa Clara until Nov. 16

Promo code: COUNTDOWN for £30 discount
for new registrations at PLUK!

Copyright 2014 Bill Karwin

www.slideshare.net/billkarwin

Released under a Creative Commons 3.0 License:
<http://creativecommons.org/licenses/by-nc-nd/3.0/>

You are free to share - to copy, distribute and transmit this work, under the following conditions:



Attribution.

You must attribute this work to Bill Karwin.

Noncommercial.

You may not use this work for commercial purposes.

No Derivative Works.

You may not alter, transform, or build upon this work.